

Some body text here to show the figure placement at the top of the page. Lorem ipsum dolor sit amet, consectetur adipiscing elit.

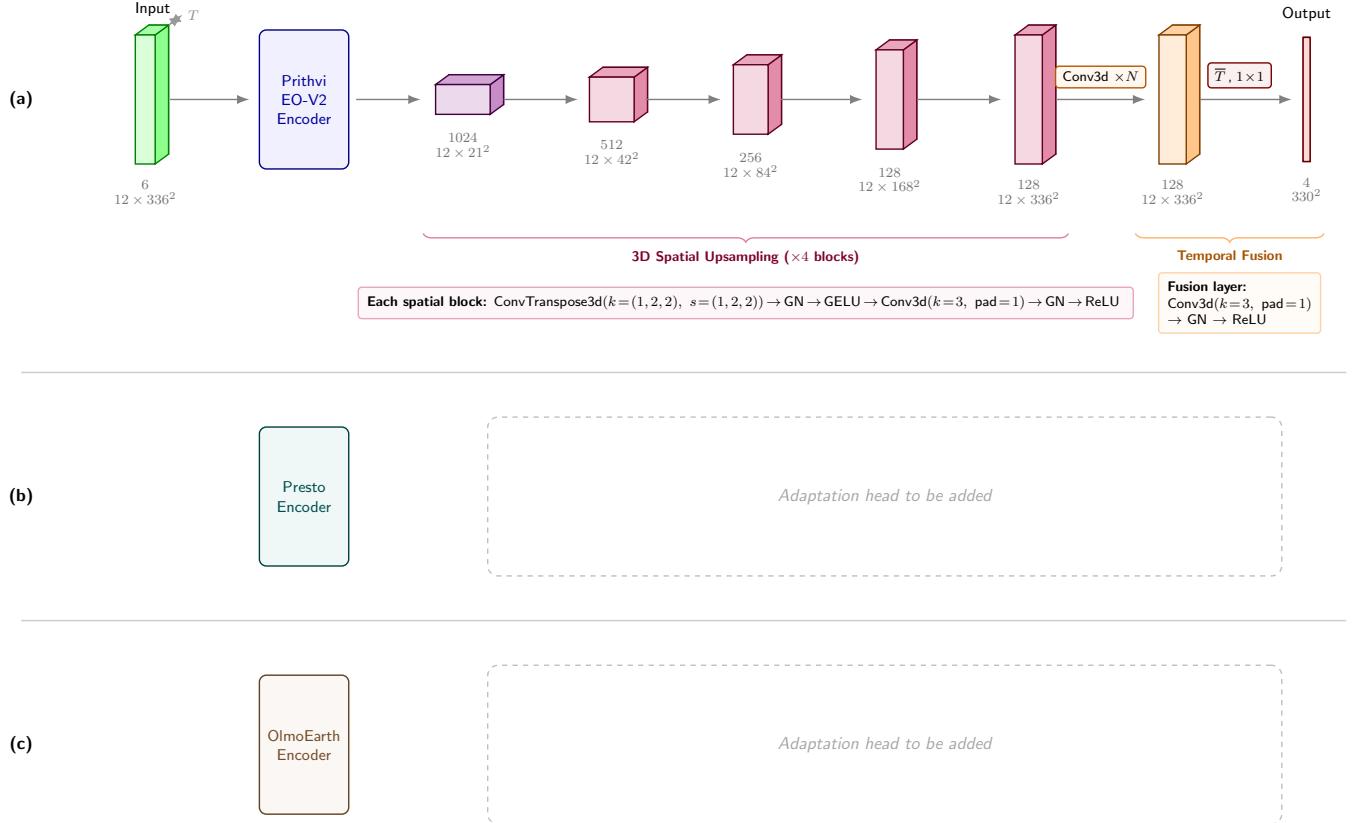


Figure 1: Downstream adaptation architectures for phenology prediction. (a) The Prithvi-EO-V2 encoder extracts spatio-temporal features, which are spatially upsampled through four 3D transposed convolution blocks while preserving the temporal dimension  $T$ . A 3D convolutional fusion head mixes temporal information, followed by mean pooling over  $T$  and a  $1 \times 1$  projection to predict four phenological transition dates per pixel. (b) Presto adaptation (to be added). (c) OlmoEarth adaptation (to be added).